

CITA: Contrastive Instruction-Tuned Alignment for Switchable LLM Behavior

Kapil Wanaskar¹ Gaytri Jena² Vaibhaav Tiwari³ Vinija Jain⁴ Aman Chadha⁵ Amitava Das⁶

¹San José State University ²GITA, India ³BML Munjal University, India

⁴Google AI ⁵Apple GenAI ⁶BITS Pilani, Goa

kapilw25@gmail.com

Abstract

Traditional LLM alignment methods like DPO embed static objectives into model weights, lacking flexibility for dynamic policy changes. We introduce **CITA (Contrastive Instruction-Tuned Alignment)**, a framework enabling **instruction-driven alignment** at inference time. CITA combines supervised fine-tuning, contrastive preference optimization, and mandatory KL regularization to achieve instruction-conditioned behavior switching.

To evaluate this capability, we contribute the **Instruction Switch Dataset (ISD)**—3,000 test cases (300 prompts $\times$ 10 instruction types) where only the alignment instruction varies while the user prompt remains constant.

Experiments on Llama-3.1-8B across 5 benchmarks show CITA outperforms DPO in 3/5 metrics (TruthfulQA, Length Control, AQI). On TruthfulQA, CITA improves by +0.054 while DPO shows minimal change (+0.001). Our results demonstrate that CITA achieves deeper instruction-alignment rather than superficial instruction-following.

Contributions at-a-glance

- **CITA Framework:** A unified loss combining SFT, DPO, and mandatory KL regularization for instruction-conditioned alignment
- **Instruction Switch Dataset (ISD):** 3,000 test cases (300 prompts $\times$ 10 instruction types) isolating alignment instruction effects
- **Comprehensive Evaluation:** CITA outperforms DPO in 3/5 benchmarks (TruthfulQA +0.054, Length Control, AQI)
- **Key Insight:** Mandatory KL regularization enables instruction-alignment (not just instruction-following)

Code: https://github.com/kapilw25/finetuning_evaluation

Dataset: <https://huggingface.co/datasets/kapilw25/ISD-Instruction-Switch-Dataset>

1 Introduction

1.1 Why Instruction-Based Alignment?

Modern AI systems increasingly operate as **autonomous agents**—from customer service chatbots to research assistants to code generation pipelines (Brown et al., 2020; OpenAI, 2023; Touvron et al., 2023b), building on foundational work in language modeling (Vaswani et al., 2017; Devlin et al., 2019; Radford et al., 2019). These agents require *dynamic behavioral adaptation* based on workflow context, user roles, and task requirements (Shinn et al., 2023). A customer-facing agent needs empathetic, safety-conscious responses; the same underlying model, when assisting researchers, should prioritize accuracy and permit nuanced discussion of sensitive topics.

The Limitation of Static Alignment. Current alignment methods like DPO (Rafailov et al., 2023), RLHF (Ouyang et al., 2022; Stiennon et al., 2020), and their variants (Azar et al., 2023; Ethayarajh et al., 2024) embed a *fixed* behavioral policy into model weights during training. Once deployed, the model cannot adapt its alignment posture—it responds identically whether serving a child, a medical professional, or a security researcher. This rigidity creates a fundamental mismatch with the flexibility demands of agentic AI systems.

Real-Time Alignment via Instructions. We propose that alignment should be **controllable at inference time** through natural language instructions. Just as AI agents receive dynamic system prompts for workflow orchestration, they should receive dynamic *alignment instructions* that modulate behavioral policies on-the-fly.

Research Questions

- **RQ1: Can alignment policies be switched at inference time?**
Can a single model exhibit different alignment behaviors based on natural language instructions, without retraining?
- **RQ2: What training modifications enable instruction-conditioned alignment?**
How does instruction-conditioned preference learning differ from standard DPO, and what regularization is required?
- **RQ3: How do we measure instruction-alignment vs. instruction-following?**
How can we distinguish deep behavioral adaptation from superficial compliance?

1.2 Motivating Scenario: AI Agent Orchestration

Consider an enterprise AI agent platform where the same base model serves multiple workflows:

Workflow Context	Required Alignment Behavior
Customer Support	Empathetic, safety-first, avoid technical jargon
Internal Research	Balanced, permit nuanced discussion, cite sources
Compliance Review	Regulatory focus, conservative, flag uncertainties
Creative Brainstorming	Exploratory, unconventional ideas welcome
Child Education	Strictly safe, educational, age-appropriate

Table 1: Same model, different alignment requirements per workflow. Static alignment cannot accommodate this flexibility.

With **static alignment** (DPO), organizations must either: (a) deploy multiple separately-aligned models, or (b) accept one-size-fits-all behavior. Neither scales. **Instruction-based alignment** enables a single model to switch behavioral modes via system prompts—the same mechanism agents already use for task orchestration.

1.3 Our Solution: CITA

We introduce **CITA (Contrastive Instruction-Tuned Alignment)**, a framework enabling **instruction-driven alignment** at inference time. CITA trains models to condition their preference behavior on alignment instructions, achieving:

- **Dynamic Policy Switching:** Different behavioral modes activated via natural language
- **No Retraining Required:** Alignment changes expressed through instructions, not weight updates
- **Agent-Compatible:** Integrates with existing system prompt mechanisms

1.4 Key Contributions

1. **CITA Framework:** A unified training objective combining SFT, contrastive preference optimization, and *mandatory* KL regularization for instruction-conditioned alignment. The key insight is that KL regularization—optional in DPO—becomes essential for stable behavioral switching.
2. **Instruction Switch Dataset (ISD):** A diagnostic benchmark of 3,000 test cases (300 prompts $\times$ 10 instruction types) designed to isolate alignment instruction effects from prompt variation. ISD enables systematic evaluation of instruction-awareness beyond surface-level instruction-following.

3. **Comprehensive Evaluation:** Experiments on Llama-3.1-8B (Dubey et al., 2024a) across 5 benchmarks demonstrate CITA outperforms DPO in 3/5 metrics. On TruthfulQA (Lin et al., 2022), CITA improves by +0.054 while DPO shows minimal change (+0.001).

1.5 Related Work

Instruction Tuning. FLAN (Wei et al., 2022; Chung et al., 2022; Longpre et al., 2023), T0 (Sanh et al., 2022), and subsequent work (Wang et al., 2023c; Iyer et al., 2022b; Mishra et al., 2022) demonstrated that instruction-following during training improves zero-shot generalization. Models like Alpaca (Taori et al., 2023), WizardLM (Xu et al., 2023), LIMA (Zhou et al., 2023), Orca (Mukherjee et al., 2023), and OpenChat (Wang et al., 2023b) further refined this paradigm, with comprehensive surveys (Peng et al., 2023) synthesizing these developments. However, these approaches focus on task generalization rather than alignment conditioning.

Preference Optimization. DPO (Rafailov et al., 2023) eliminates the need for separate reward models, inspiring variants like KTO (Ethayarajh et al., 2024), ORPO (Hong et al., 2024), SimPO (Meng et al., 2024), and RRHF (Yuan et al., 2023). However, all these methods lack instruction-awareness—they embed static alignment into weights. Recent work (Xu et al., 2024) compares DPO and PPO (Schulman et al., 2017) but neither supports dynamic policy switching.

Safety Alignment. Constitutional AI (Bai et al., 2022b), PKU-SafeRLHF (Ji et al., 2024), and red-teaming approaches (Perez et al., 2022; Ganguli et al., 2022) focus on making models safer but assume a single, fixed safety policy. Research on jailbreaking (Zou et al., 2023; Wei et al., 2023; Liu et al., 2023) and fine-tuning risks (Qi et al., 2023; Huang et al., 2023) reveals vulnerabilities in static alignment. Recent work on context-dependent safety (Bianchi et al., 2024) highlights the need for adaptive policies.

Agentic AI Systems. LLM-based agents (Wang et al., 2024; Xi et al., 2023) increasingly rely on dynamic system prompts for workflow orchestration. CITA extends this paradigm to alignment itself, enabling behavioral policy changes through the same mechanism.

Property	DPO	CITA
Reward Model Required	No	No
Instruction-Aware	No	Yes
Behavioral Switching	No	Yes
Mandatory KL	Optional	Yes
Dynamic Policy	No	Yes
Agent-Compatible	No	Yes

Table 2: Comparison of alignment methods. CITA uniquely enables instruction-conditioned behavioral switching compatible with agentic AI workflows.

Positioning. CITA bridges instruction tuning and preference optimization by training models to condition their alignment behavior on natural language instructions. This enables a new paradigm of **interactive alignment** where policy changes can be expressed without retraining—essential for the flexibility demands of modern AI agent systems.

2 Methodology

We present the CITA training methodology, organized into three phases: (1) setup and data preparation, (2) instruction-tuned formatting, and (3) contrastive training configuration.

2.1 Phase 0: Model and Infrastructure Setup

Component	Specification
Base Model	Llama-3.1-8B (Dubey et al., 2024a)
Prompt Format	Chat-style (system/user/assistant)
Tokenizer	LLaMA-3 with BOS/EOS tokens
Hardware	NVIDIA A100-40GB
Training Data	PKU-SafeRLHF (Ji et al., 2024)

Table 3: Training infrastructure and configuration.

Rationale. We select Llama-3.1-8B for its strong instruction-following capabilities and accessibility, following the Llama model family tradition (Touvron et al., 2023b,a; Dubey et al., 2024b). PKU-SafeRLHF provides high-quality preference pairs with explicit safety annotations (Bai et al., 2022a), enabling evaluation of instruction-conditioned safety behaviors. Alternative datasets like OpenAssistant (Köpf et al., 2023), UltraFeedback (Cui et al., 2023), and Tulu (Iverson et al., 2023) could also be used.

2.2 Phase 1: Dataset Formatting

We define two complementary data formats that enable instruction-conditioned training:

Instruction-Tuned Alignment (ITA) Format. Each training example explicitly includes an alignment instruction alongside the user query:

```

<user>
### Alignment Instruction:
${alignment_instruction}
### User Prompt:
${user_prompt}
<assistant>
${expected_response}

```

Contrastive ITA Format. For preference learning, we extend ITA to include rejected responses:

```

### Rejected Response (Bad):
${rejected_response}
### Expected Response (Good):
${accepted_response}

```

Example-Based Alignment (EBA) Format. Standard DPO format without explicit instructions for baseline comparison:

```
{prompt, chosen, rejected}
```

2.3 Phase 2: Training Configuration

Parameter	SFT Stage	CITA Stage
Trainer	SFTTrainer	Custom DPO+KL
Epochs	3–5	3
Seq. Length	2048	2048
Learning Rate	2e-5	5–7e-6
LoRA Rank	8	8
β (temp.)	—	0.1–0.12
λ_{KL}	—	0.0002–0.0005

Table 4: Training hyperparameters for SFT and CITA stages.

Key Design Choice. The CITA stage uses significantly lower learning rates ($3\text{--}4\times$ lower than typical DPO) because instruction-augmented sequences are 30–40% longer, producing larger gradients. This finding emerged from hyperparameter optimization using Optuna (Akiba et al., 2019) (Section 6.4). We use parameter-efficient fine-tuning via LoRA (Hu et al., 2022; Dettmers et al., 2023) rather than full fine-tuning to enable single-GPU training. Alternative PEFT methods include adapter tuning (Houlsby et al., 2019), prompt tuning (Lester et al., 2021), and prefix tuning (Li and Liang, 2021).

3 Unified Training Objective

CITA combines three complementary loss components to achieve instruction-conditioned alignment, drawing on insights from supervised fine-tuning (Wei et al., 2022; Sanh et al., 2022), contrastive preference learning (Rafailov et al., 2023; Chen et al., 2020), and KL-regularized optimization (Schulman et al., 2017; Korbak et al., 2022). The key insight is that each component serves a distinct purpose, and their combination—particularly

the *mandatory* KL term—enables stable behavioral switching.

3.1 Loss Formulation

The unified CITA objective is:

$$\mathcal{L}_{\text{CITA}} = \mathcal{L}_{\text{SFT}} + \lambda_1 \mathcal{L}_{\text{DPO}} + \lambda_2 \mathcal{L}_{\text{KL}} \quad (1)$$

where $\lambda_1, \lambda_2 > 0$ are balancing coefficients. We now detail each component:

SFT Loss (Response Quality). Standard cross-entropy loss ensuring high-quality response generation:

$$\mathcal{L}_{\text{SFT}} = -\frac{1}{N} \sum_{(x, y^+)} \sum_{t=1}^T \log P_{\theta}(y_t^+ | x, y_{<t}^+) \quad (2)$$

where x is the instruction-augmented prompt, y^+ is the preferred response, and T is the sequence length.

DPO Loss (Preference Alignment). Contrastive preference optimization without a separate reward model (Rafailov et al., 2023; Ziegler et al., 2019):

$$\mathcal{L}_{\text{DPO}} = -\frac{1}{N} \sum \log \sigma(\beta \cdot \Delta_{\theta}) \quad (3)$$

where the log-probability difference is:

$$\Delta_{\theta} = \log P_{\theta}(y^+ | x) - \log P_{\theta}(y^- | x) \quad (4)$$

and $\beta > 0$ is the temperature parameter controlling preference sharpness.

KL Loss (Stability — Mandatory). The critical component enabling behavioral switching, inspired by KL-regularized RL (Jaques et al., 2019; Korbak et al., 2022) and continual learning (Kirkpatrick et al., 2017; Li and Hoiem, 2017):

$$\mathcal{L}_{\text{KL}} = \frac{1}{N} \sum_{(x, y)} \text{KL}[P_{\theta}(\cdot | x) \| P_{\text{ref}}(\cdot | x)] \quad (5)$$

where P_{ref} is the reference model (pre-CITA checkpoint).

3.2 Why Mandatory KL?

Scenario	Without KL	With KL
Mode Collapse	High risk	Prevented
Instruction Switching	Unstable	Stable
Behavioral Consistency	Degrades	Maintained
Preference Margin	Unbounded	Controlled

Table 5: Effect of mandatory KL regularization on training stability.

Key Insight. In standard DPO, KL regularization is implicit and optional. For CITA, we make it *explicit and mandatory* because:

1. **Prevents Mode Collapse:** Without KL, the model can overfit to specific instruction patterns, losing generalization—a form of catastrophic forgetting (French, 1999; Kirkpatrick et al., 2017).
2. **Enables Switching:** KL keeps the model “close enough” to the base policy to respond appropriately to different instructions rather than collapsing to a single behavioral mode, similar to knowledge distillation (Hinton et al., 2015).
3. **Bounds Preference Margins:** Unconstrained optimization can lead to extreme log-probability differences, destabilizing training—a known issue in PPO (Schulman et al., 2017), DPO variants (Azar et al., 2023; Zhao et al., 2023), and gradient-based continual learning (Lopez-Paz and Ranzato, 2017).

3.3 Gradient Analysis

The gradient of the contrastive term with respect to model parameters is:

$$\nabla_{\theta} \mathcal{L}_{\text{DPO}} = -\beta [(1 - P^+) \nabla_{\theta} s^+ - P^+ \nabla_{\theta} s^-] \quad (6)$$

where $P^+ = \sigma(\beta \cdot \Delta_{\theta})$ is the preference probability and $s^{\pm} = \log P_{\theta}(y^{\pm} | x)$.

Interpretation. When $P^+ \approx 1$ (model correctly prefers y^+), gradients become small, preventing over-optimization. The KL term adds a stabilizing force that prevents P^+ from saturating too quickly.

4 CITA: Contrastive Instruction-Tuned Alignment

Building on the unified loss formulation (Section 3), we present the complete CITA framework, including the training triplet structure, the full CITA-KL loss, and the training pipeline.

4.1 Training Triplet

Each CITA training example consists of a 4-tuple:

$$(I, X, Y^+, Y^-) \in \mathcal{D}_{\text{CITA}} \quad (7)$$

where:

- I : Alignment instruction (e.g., “Respond with safety-first principles”)

- X : User input/query
- Y^+ : Preferred response (aligned with instruction I)
- Y^- : Rejected response (misaligned with instruction I)

Crucial Difference from DPO. In standard DPO (Rafailov et al., 2023), the instruction I is absent or implicit. CITA *explicitly conditions* preference on I , enabling the model to learn instruction-specific behavioral patterns—analogueous to how contrastive learning (Chen et al., 2020; Khosla et al., 2020) conditions representations on class labels.

4.2 CITA-KL Loss

The complete CITA objective combines contrastive learning (Oord et al., 2018; Gao et al., 2021; Radford et al., 2021) with mandatory KL (Korbak et al., 2022):

$$\mathcal{L}_{\text{CITA-KL}} = -\log P^+ + \lambda \cdot \text{KL}[\pi_{\theta} \| \pi_0] \quad (8)$$

where the preference probability is:

$$P^+ = \frac{\exp(\beta \cdot s^+)}{\exp(\beta \cdot s^+) + \exp(\beta \cdot s^-)} \quad (9)$$

with instruction-conditioned scores:

$$s^+ = \log \pi_{\theta}(Y^+ | I, X) \quad (10)$$

$$s^- = \log \pi_{\theta}(Y^- | I, X) \quad (11)$$

4.3 Hyperparameters

Symbol	Description
π_{θ}	Fine-tuned CITA model
π_0	Reference model (pre-CITA checkpoint)
β	Temperature (controls preference sharpness)
λ	KL weight (mandatory , > 0)

Table 6: CITA-KL hyperparameters. The KL weight λ is non-zero by design.

4.4 Training Pipeline

CITA follows a staged training approach, where each stage builds on the previous checkpoint:




Figure 1: CITA training pipeline. Each stage uses the previous checkpoint. CITA adds mandatory KL regularization for instruction-conditioned alignment.

4.5 Benefits of CITA

1. **Generalization:** Responds appropriately to novel instruction-prompt combinations not seen during training, similar to zero-shot generalization in instruction tuning (Wei et al., 2022; Chung et al., 2022).
2. **Robustness:** Resists jailbreak attempts (Zou et al., 2023; Wei et al., 2023) by conditioning behavior on explicit alignment instructions.
3. **Calibration:** Avoids over-refusal (Bianchi et al., 2023) by adjusting safety level based on instruction context.
4. **Behavioral Switching:** Dynamically adapts response style (e.g., concise vs. detailed, formal vs. casual) based on instructions, enabling multi-tenant deployment (Wang et al., 2024).

4.6 CITA vs. Prior Methods

Capability	DPO	CITA
No Reward Model	✓	✓
Instruction-Conditioned	✗	✓
Dynamic Policy Switch	✗	✓
Explicit KL Control	Implicit	Mandatory

Table 7: Feature comparison across alignment methods. CITA uniquely supports instruction-conditioned behavioral switching.

5 Instruction Switch Dataset (ISD)

We introduce the **Instruction Switch Dataset (ISD)**—a diagnostic benchmark specifically designed to evaluate instruction-conditioned alignment. Unlike existing benchmarks (Zheng et al., 2023; Wang et al., 2023a; Liang et al., 2022) that conflate instruction effects with prompt variation, ISD isolates the *pure effect* of alignment instructions by keeping prompts constant.

Dataset Design Questions

- **DQ1: How do we isolate instruction effects from prompt effects?**
Keep user prompts constant while varying only the alignment instruction.
- **DQ2: What instruction types capture meaningful behavioral dimensions?**
Select 10 types spanning perspective, style, and constraint dimensions.

5.1 Motivation

Existing benchmarks like TruthfulQA (Lin et al., 2022), MMLU (Hendrycks et al., 2021b), Hel-laSwag (Zellers et al., 2019), and ARC (Clark et al., 2018) conflate instruction-following with prompt variation. When both change simultaneously, it is impossible to determine whether behavioral differences stem from instruction comprehension or prompt-specific responses.

ISD Design Principle. By fixing the user prompt X and varying only the instruction I , we can measure:

$$\Delta_I = \mathbb{E}_{I_1, I_2} [d(R(I_1, X), R(I_2, X))] \quad (12)$$

where $R(I, X)$ is the response to prompt X under instruction I , and $d(\cdot, \cdot)$ measures behavioral difference.

5.2 Dataset Construction

Prompt Bank. We curate 500 prompts from diverse sources:

- News articles and current events
- Educational Q&A collections
- Instruction-tuning datasets (FLAN (Wei et al., 2022), Alpaca (Taori et al., 2023), OPT-IML (Iyer et al., 2022a))
- User intent pools from dialogue systems

Selection Criteria. All prompts are:

- **Semantically open-ended:** Allow multiple valid response styles
- **Culturally grounded:** Avoid region-specific assumptions
- **Neutrally phrased:** No embedded biases toward specific instruction types

5.3 10 Instruction Types

Instruction Type	Behavioral Objective
neutral	Balanced perspective, present pros/cons
conservative	Favor traditional methods, risk-averse
liberal	Encourage innovation, embrace change
regulatory	Emphasize compliance, legal considerations
empathetic	Compassionate, emotionally supportive tone
safety_first	Prioritize risk awareness, caution
educational	Teaching mode, explain concepts clearly
concise	Brief, direct responses
professional	Formal tone, business-appropriate
creative	Novel approaches, unconventional ideas

Table 8: The 10 instruction types in ISD, each targeting a distinct behavioral dimension.

5.4 Dataset Scale

The ISD benchmark is constructed as a full factorial design:

$$|\text{ISD}| = 300 \times 10 = \mathbf{3,000} \text{ test cases} \quad (13)$$

This structure enables systematic evaluation: each prompt appears exactly 10 times (once per instruction type), allowing direct comparison of behavioral responses while controlling for prompt-specific effects. The 300 base prompts are drawn from 12 categories (25 prompts each), ensuring topical diversity across economics, healthcare, technology, ethics, and other domains.

5.5 Example: Instruction Switching

The core value of ISD is demonstrating how a **single prompt** produces **10 distinct behavioral responses** based solely on the alignment instruction. Table 9 shows this for a representative prompt.

Prompt: “Should AI-generated art be allowed in competitions?”

Instruction	Expected Behavioral Response
neutral	Present pros and cons without taking sides; acknowledge both innovation benefits and authenticity concerns.
conservative	Favor traditional human craftsmanship; competitions should reward proven artistic skills and techniques.
liberal	Embrace innovation and democratization; AI expands creative boundaries and enables new artistic voices.
regulatory	Discuss copyright implications, disclosure requirements, and labeling standards for AI-generated entries.
empathetic	Acknowledge artists’ livelihood concerns while recognizing the excitement of new creative tools.
safety_first	Warn about authenticity deception risks, potential for fraud, and undermining trust in art institutions.
educational	Explain how AI art generation works, its capabilities, limitations, and implications for art education.
concise	Brief direct answer: “Yes, with disclosure requirements” or “No, separate category needed.”
professional	Formal industry analysis discussing market implications, institutional policies, and stakeholder interests.
creative	Explore unconventional possibilities: AI-human collaborations, new competition categories, hybrid art forms.

Table 9: **Instruction Switching Example:** Same prompt produces 10 behaviorally distinct responses. Each instruction activates a different alignment policy, demonstrating CITA’s dynamic behavioral switching capability.

5.6 Prompt Category Distribution

ISD prompts span 12 diverse categories to ensure broad coverage of real-world topics:

Category	#	Category	#
Technology	25	Ethics	25
Healthcare	25	Culture	25
Environment	25	Science	25
Education	25	Business	25
Economics	25	Personal	25
Social	25	Governance	25

Table 10: ISD prompt distribution across 12 categories (300 prompts total).

5.7 Instruction-Following vs. Instruction-Alignment

A critical distinction underlies the ISD evaluation:

Property	Following	Alignment
Depth	Surface-level	Internalized
Example	“Be concise” → short text	Concise <i>style</i> adopted
Robustness	Easily disrupted	Consistent across contexts
Goal	Format compliance	Policy embodiment

Table 11: Instruction-following (shallow) vs. instruction-alignment (deep).

ISD tests whether CITA achieves true instruction-alignment—consistent behavioral shifts reflecting internalized policy goals—rather than superficial instruction-following that can be easily circumvented, building on recent work in self-play fine-tuning (Chen et al., 2024).

5.8 Dataset Availability

The Instruction Switch Dataset is publicly available:

<https://huggingface.co/datasets/kapilw25/ISD-Instruction-Switch-Dataset>

6 Experiments

We present our experimental setup, including the training pipeline, model variants, evaluation benchmarks, hyperparameter optimization, and training dynamics.

6.1 Training Pipeline

Following the staged approach from Section 4.4, we train on PKU-SafeRLHF (Ji et al., 2024; Bai et al., 2022a) using an NVIDIA A100-40GB GPU with LoRA (Hu et al., 2022) adapters. Training proceeds through four stages:

1. **Base:** Llama-3.1-8B (Dubey et al., 2024a) (pre-trained weights)

2. **SFT**: Fine-tune on PKU chosen responses
3. **DPO**: Train on PKU preference pairs
4. **CITA**: Add instruction-conditioning with mandatory KL

6.2 Model Variants

We train 6 model variants across 3 methods (SFT, DPO, CITA) $\times$ 2 instruction settings (NoInstruct, Instruct):

Model	Training Path	Instruction
SFT_NI	Base→SFT	No
SFT_I	Base→SFT	Yes
DPO_NI	Base→SFT→DPO	No
DPO_I	Base→SFT→DPO	Yes
CITA_NI	Base→SFT→DPO→CITA	No
CITA_I	Base→SFT→DPO→CITA	Yes

Table 12: 6 model variants. NI = NoInstruct (standard prompts), I = Instruct (with behavioral instructions).

Comparison Design. By evaluating NoInstruct vs. Instruct variants, we measure each method’s *instruction sensitivity*—the improvement from adding alignment instructions.

6.3 Evaluation Benchmarks

We evaluate on 5 benchmarks with deterministic metrics (no LLM-as-judge):

Benchmark	Cases	Instr. Types
ISD (ours)	3,000	10 types
TruthfulQA	1,634	HON/CONF
Cond. Safety	1,000	STRICT/PERMIS.
Length Control	1,000	CONC./DETAIL.
AQI	2,800	7 axioms

Table 13: Evaluation benchmarks. All use deterministic metrics.

Benchmark	Metric (Ideal Value)
ISD	Fidelity $\times$ Behavioral Shift (1.0)
TruthfulQA	HON – CONF uncertainty difference (+)
Conditional Safety	STRICT – PERMISSIVE (1.0)
Length Control	DETAILED/CONCISE word ratio (>4)
AQI	(CHI + XB) / 2 clustering score (100)

Table 14: Metric definitions. Higher is better for all except TruthfulQA (positive direction matters).

Metrics. **ISD Types:** Neutral, Conservative, Liberal, Regulatory, Empathetic, Safety, Educational, Concise, Professional, Creative.

AQI Axioms: Civility, Duty, Empathy, Information, Justice, Well-being, Wisdom.

6.4 Hyperparameter Optimization

We use Optuna (Akiba et al., 2019) with TPE sampler (12 trials) to optimize hyperparameters, following best practices from hyperparameter optimization literature (Liang et al., 2022). Key finding: **CITA_Instruct requires $\sim 50\%$ lower learning rate** than CITA_NoInstruct because instruction-augmented sequences are 30–40% longer, causing larger gradients.

Hyperparameter	NoInstruct (T5)	Instruct (T7)
Learning rate	6.83e-6	5.41e-6
λ_{KL}	0.00052	0.00023
β (DPO temp.)	0.119	0.107
Weight decay	0.0091	0.0109
Warmup ratio	7.5%	10.0%
Final margin	6.95	7.52
Accuracy	89.5%	89.0%

Table 15: Optimal hyperparameters from Optuna search. T5/T7 = trial number of best configuration.

6.5 Training Dynamics

We analyze key training curves comparing DPO and CITA variants. Additional training curves (eval loss, SFT loss, SFT token accuracy) are provided in Appendix I.




Figure 2: **Reward Accuracy:** All models converge to 89–92%. DPO variants achieve slightly higher accuracy.




Figure 3: **Reward Margins:** CITA_Instruct (7.5) $>$ CITA_NoInstruct (7.1) $>$ DPO (~ 6.0). Higher margins indicate stronger preference learning.

Key Observation. CITA achieves higher reward margins (7.5 vs. 6.0) than DPO despite similar accuracy, validating that mandatory KL regularization leads to stronger preference differentiation.

7 Results

We present evaluation results across 5 benchmarks, comparing 6 model variants. We use deterministic metrics following best practices from holistic evaluation (Liang et al., 2022; Sun et al., 2024). The key finding is that **CITA outperforms DPO in 3/5 metrics** when measuring instruction-sensitivity (NoInstruct $\rightarrow$ Instruct improvement).

7.1 Overall Performance

Model	ISD	TQA	CS	LC	AQI
SFT_NI	.126	-.006	.002	0.95	37.6
SFT_I	.142	.012	.022	1.02	47.7
DPO_NI	.217	-.006	.015	0.99	11.0
DPO_I	.389	-.005	.489	1.12	55.0
CITA_NI	.204	-.040	.009	0.98	9.4
CITA_I	.367	.013	.400	1.14	55.5

Table 16: Results across 5 benchmarks. TQA = TruthfulQA, CS = Conditional Safety, LC = Length Control. **Bold** = best per column. NI = NoInstruct, I = Instruct.

7.2 Instruction Sensitivity: CITA vs. DPO

The critical comparison is the *improvement* from NoInstruct to Instruct, measuring how well each method responds to alignment instructions:

Benchmark	SFT Δ	DPO Δ	CITA Δ	Winner
ISD	+0.016	+0.172	+0.162	DPO
TruthfulQA	+0.018	+0.001	+0.054	CITA
Cond. Safety	+0.020	+0.475	+0.391	DPO
Length Control	+0.063	+0.130	+0.164	CITA
AQI	+10.1	+44.0	+46.1	CITA

Table 17: Instruction sensitivity (Δ = Instruct – NoInstruct). **CITA wins 3/5, DPO wins 2/5, SFT wins 0/5.**

Final Score: CITA 3 – DPO 2 – SFT 0

7.3 Benchmark-Specific Analysis




Figure 4: **ISD Results:** DPO_Instruct (0.389) and CITA_Instruct (0.367) show highest instruction awareness.




Figure 5: **ISD Lollipop:** Instruction sensitivity scores with emphasis on relative differences.




Figure 6: **TruthfulQA:** Only CITA_Instruct (+0.013) and SFT_Instruct (+0.012) achieve *positive* adaptation scores.




Figure 7: **TruthfulQA Lollipop:** Adaptation scores highlighting CITA's superior uncertainty calibration.




Figure 8: **Conditional Safety:** DPO_Instruct (0.489) and CITA_Instruct (0.400) show behavioral gap between STRICT vs. PERMISSIVE.




Figure 9: **Conditional Safety Lollipop**: Safety adaptation scores across STRICT vs. PERMISSIVE contexts.




Figure 10: **Length Control**: CITA_Instruct valid-only (1.77) shows strongest adaptation.




Figure 11: **Length Control Lollipop**: Response length adaptation ratios with baseline reference.




Figure 12: **AQI**: CITA_Instruct (55.5) and DPO_Instruct (55.0) achieve highest clustering scores.




Figure 13: **AQI Lollipop**: Alignment quality scores emphasizing CITA and DPO performance gap.

7.4 Combined Analysis




Figure 15: **Performance Heatmap**: Original scores with column-normalized colors. Green = best, red = worst.

7.5 Key Insights

TruthfulQA is the Differentiator. On TruthfulQA, CITA improves by +0.054 (becoming appropriately more/less confident when instructed), while DPO shows negligible change (+0.001). This demonstrates CITA’s superior ability to *internalize* uncertainty calibration instructions.

CITA Achieves Higher Margins with Lower Accuracy. Despite similar training accuracy (~89%), CITA achieves higher reward margins (7.5 vs. 6.0). This validates that mandatory KL prevents margin collapse while enabling stable behavioral switching.

Instruction-Alignment $\neq$ Instruction-Following. DPO_I achieves higher absolute ISD scores, but CITA_I shows more consistent improvement across diverse benchmarks. This suggests CITA achieves *deeper* instruction-alignment rather than surface-level following.

8 Conclusion

This paper introduces **CITA (Contrastive Instruction-Tuned Alignment)**, a framework enabling instruction-aware alignment in large language models. Unlike DPO (Rafailov et al.,




Figure 14: **Instruction Adaptation Radar**: CITA wins 3/5 benchmarks, DPO wins 2/5, SFT wins 0/5. Values show NoInstruct → Instruct improvement.

2023) and RLHF (Ouyang et al., 2022) which embed static alignment into weights, CITA allows models to dynamically adapt behavior based on natural language instructions at inference time.

8.1 Key Findings

- CITA outperforms DPO in instruction-awareness**: Measuring NoInstruct → Instruct improvement, CITA wins **3/5 benchmarks** (TruthfulQA, Length Control, AQI), while DPO wins 2/5 (ISD, Conditional Safety).
- Correct directional adaptation**: On TruthfulQA, CITA improves by **+0.054** (adjusting uncertainty appropriately), while DPO shows minimal change (**+0.001**). This validates CITA’s ability to internalize calibration instructions.
- Mandatory KL is essential**: The KL term—optional in DPO—becomes mandatory in CITA to prevent mode collapse and enable stable instruction-conditioned behavioral switching.
- Instruction-alignment $\neq$ instruction-following**: Through the Instruction Switch Dataset, we demonstrate that CITA achieves *deeper* instruction-alignment (internalized policies) rather than *superficial* instruction-following (surface compliance).

8.2 Implications

CITA enables a new paradigm of **interactive and adaptive AI systems** where:

- Policy updates via natural language**: Alignment changes can be expressed through instructions without model retraining.
- Context-aware responses**: Different users, applications, or deployment contexts can receive appropriately calibrated behaviors.
- Dynamic safety policies**: Safety constraints can be adjusted at inference time based on deployment requirements (e.g., stricter for child-facing apps, more permissive for research contexts).
- Reduced alignment tax**: No need for multiple specialized models for different alignment policies—a single CITA model can switch behaviors on demand.

Outlook. CITA opens several research directions: (1) scaling to larger models (Touvron et al., 2023b; Jiang et al., 2023; Team et al., 2023; Chowdhery et al., 2022) and more instruction types, (2) theoretical analysis of instruction-conditioned learning dynamics (Azar et al., 2023), (3) applications to multi-stakeholder alignment where differ-

ent parties have different policy requirements (Sun et al., 2024), and (4) integration with constitutional AI methods (Bai et al., 2022b) for hierarchical instruction handling.

8.3 Reproducibility

All code, trained models, and the Instruction Switch Dataset are publicly available:

- **Code:** https://github.com/kapilw25/finetuning_evaluation
- **Dataset:** <https://huggingface.co/datasets/kapilw25/ISD-Instruction-Switch-Dataset>

9 Limitations

We acknowledge the current limitations of CITA and outline directions for future work.

9.1 Experimental Scope

Limitation	Impact & Mitigation Path
Single model (Llama-3.1-8B)	Validate on GPT, Mistral, Gemma; scale to 70B
English-only evaluation	Extend to multilingual instruction transfer
10 instruction types	Add domain-specific types (medical, legal)
Single dataset (PKU-SafeRLHF)	Test on diverse preference datasets

Table 18: Current experimental limitations and mitigation paths.

Model Generalization. We evaluate only Llama-3.1-8B (Dubey et al., 2024a). Future work should validate CITA on diverse architectures (Jiang et al., 2023; Zhang et al., 2022; OpenAI, 2023) and scales (7B–70B) to establish generality across model families.

Instruction Coverage. The 10 instruction types in ISD cover common alignment dimensions but may not capture all real-world policy requirements. Domain-specific instructions (medical, legal, educational) require dedicated evaluation.

Multilingual Alignment. All experiments are English-only. Instruction-conditioned alignment in multilingual settings raises questions about cross-lingual instruction transfer and cultural alignment differences.

9.2 Deployment Considerations

Inference Overhead. CITA requires alignment instructions in every prompt, increasing context length by 10–40 tokens. While minimal for modern LLMs, this may matter for latency-critical applications or extremely long contexts.

Instruction Specification. Deployers must carefully craft alignment instructions for each use case. Poorly specified instructions may lead to unexpected behaviors—instruction engineering becomes a new deployment concern.

9.3 Safety & Ethical Considerations

Dual-Use Risk. Instruction-conditioned alignment could be misused to make models *less* safe via adversarial instructions (Zou et al., 2023; Carlini et al., 2023; Pace et al., 2024). Mitigations include:

- Hierarchical instruction handling (system > user)
- Instruction validation/filtering at deployment
- Constitutional constraints (Bai et al., 2022b) as non-negotiable baselines

Alignment Washing. Organizations might claim “aligned” behavior while using permissive instructions. Transparency about instruction policies and third-party auditing are essential safeguards.

Control Boundaries. CITA enables user-influenced alignment, raising questions about who controls behavioral policy. We advocate for tiered systems: users adjust within deployer-set bounds, which operate within regulatory constraints.

9.4 Future Work

1. **Hierarchical Instructions:** Multi-level handling where system policies override user preferences
2. **Instruction Compositionality:** Combining multiple instructions (“be concise AND professional”) meaningfully
3. **Instruction Robustness:** Resistance to instruction injection attacks
4. **Theoretical Analysis:** Convergence guarantees and sample complexity for instruction-conditioned learning
5. **Multi-Modal CITA:** Extension to vision-language models with cross-modal instructions

References

- Takuya Akiba, Shotaro Sano, Toshihiko Yanase, Takeru Ohta, and Masanori Koyama. 2019. [Optuna: A next-generation hyperparameter optimization framework](#). *Proceedings of the 25th ACM SIGKDD International Conference on Knowledge Discovery & Data Mining*, pages 2623–2631.
- Mohammad Gheshlaghi Azar, Mark Rowland, Bilal Piot, Daniel Guo, Daniele Calandriello, Michal Valko, and Rémi Munos. 2023. [A general theoretical paradigm to understand learning from human preferences](#). *arXiv preprint arXiv:2310.12036*.
- Yuntao Bai, Andy Jones, Kamal Ndousse, Amanda Askell, Anna Chen, Nova DasSarma, et al. 2022a. [Training a helpful and harmless assistant with reinforcement learning from human feedback](#). *arXiv preprint arXiv:2204.05862*.
- Yuntao Bai, Saurav Kadavath, Sandipan Kundu, Amanda Askell, Jackson Kernion, Andy Jones, Anna Chen, et al. 2022b. [Constitutional ai: Harmlessness from ai feedback](#). *arXiv preprint arXiv:2212.08073*.
- Federico Bianchi, Mirac Suzgun, Giuseppe Attanasio, Paul Röttger, Dan Jurafsky, Tatsunori Hashimoto, and James Zou. 2023. [Safety-tuned llamas: Lessons from improving the safety of large language models that follow instructions](#). *arXiv preprint arXiv:2309.07875*.
- Federico Bianchi, Mirac Suzgun, Giuseppe Attanasio, Paul Röttger, Dan Jurafsky, Tatsunori Hashimoto, and James Zou. 2024. [Safety-tuned llamas: Lessons from improving the safety of large language models that follow instructions](#). *International Conference on Learning Representations*.
- Tom Brown, Benjamin Mann, Nick Ryder, Melanie Subbiah, Jared D Kaplan, Prafulla Dhariwal, et al. 2020. [Language models are few-shot learners](#). *Advances in Neural Information Processing Systems*, 33:1877–1901.
- Nicholas Carlini, Milad Nasr, Christopher A Choquette-Choo, Matthew Jagielski, Irena Gao, Anas Awadalla, et al. 2023. [Are aligned neural networks adversarially aligned?](#) *arXiv preprint arXiv:2306.15447*.
- Ting Chen, Simon Kornblith, Mohammad Norouzi, and Geoffrey Hinton. 2020. [A simple framework for contrastive learning of visual representations](#). *International Conference on Machine Learning*, pages 1597–1607.
- Zixiang Chen, Yihe Deng, Huizhuo Yuan, Kaixuan Ji, and Quanquan Gu. 2024. [Self-play fine-tuning converts weak language models to strong language models](#). *arXiv preprint arXiv:2401.01335*.
- Aakanksha Chowdhery, Sharan Narang, Jacob Devlin, Maarten Bosma, Gaurav Mishra, Adam Roberts, et al. 2022. [Palm: Scaling language modeling with pathways](#). *arXiv preprint arXiv:2204.02311*.
- Hyung Won Chung, Le Hou, Shayne Longpre, Barret Zoph, Yi Tay, William Fedus, Yunxuan Li, Xuezhi Wang, et al. 2022. [Scaling instruction-finetuned language models](#). *arXiv preprint arXiv:2210.11416*.
- Peter Clark, Isaac Cowhey, Oren Etzioni, Tushar Khot, Ashish Sabharwal, Carissa Schoenick, and Oyvind Tafjord. 2018. [Think you have solved question answering? try arc, the ai2 reasoning challenge](#). *arXiv preprint arXiv:1803.05457*.
- Ganqu Cui, Lifan Yuan, Ning Ding, Guanming Yao, Wei Zhu, Yuan Ni, et al. 2023. [Ultrafeedback: Boosting language models with high-quality feedback](#). *arXiv preprint arXiv:2310.01377*.
- Tim Dettmers, Artidoro Pagnoni, Ari Holtzman, and Luke Zettlemoyer. 2023. [Qlora: Efficient finetuning of quantized llms](#). *Advances in Neural Information Processing Systems*.
- Jacob Devlin, Ming-Wei Chang, Kenton Lee, and Kristina Toutanova. 2019. [Bert: Pre-training of deep bidirectional transformers for language understanding](#). *Proceedings of NAACL-HLT*, pages 4171–4186.
- Abhimanyu Dubey, Abhinav Jauhri, Abhinav Pandey, Abhishek Kadian, Ahmad Al-Dahle, Aiesha Letman, Akhil Mathur, Alan Schelten, Amy Yang, Angela Fan, et al. 2024a. [The llama 3 herd of models](#). *arXiv preprint arXiv:2407.21783*.
- Abhimanyu Dubey, Abhinav Jauhri, Abhinav Pandey, Abhishek Kadian, Ahmad Al-Dahle, et al. 2024b. [The llama 3 herd of models](#). *arXiv preprint arXiv:2407.21783*.
- Kawin Ethayarajh, Winnie Xu, Niklas Muennighoff, Dan Jurafsky, and Douwe Kiela. 2024. [Kto: Model alignment as prospect theoretic optimization](#). *arXiv preprint arXiv:2402.01306*.
- Robert M French. 1999. [Catastrophic forgetting in connectionist networks](#). *Trends in Cognitive Sciences*, 3(4):128–135.
- Deep Ganguli, Liane Lovitt, Jackson Kernion, Amanda Askell, Yuntao Bai, Saurav Kadavath, et al. 2022. [Red teaming language models to reduce harms: Methods, scaling behaviors, and lessons learned](#). *arXiv preprint arXiv:2209.07858*.
- Tianyu Gao, Xingcheng Yao, and Danqi Chen. 2021. [Simcse: Simple contrastive learning of sentence embeddings](#). *Proceedings of EMNLP*, pages 6894–6910.
- Dan Hendrycks, Collin Burns, Steven Basart, Andrew Critch, Jerry Li, Dawn Song, and Jacob Steinhardt. 2021a. [Aligning ai with shared human values](#). *International Conference on Learning Representations*.
- Dan Hendrycks, Collin Burns, Steven Basart, Andy Zou, Mantas Mazeika, Dawn Song, and Jacob Steinhardt. 2021b. [Measuring massive multitask language understanding](#). *International Conference on Learning Representations*.

- Geoffrey Hinton, Oriol Vinyals, and Jeff Dean. 2015. [Distilling the knowledge in a neural network](#). *arXiv preprint arXiv:1503.02531*.
- Jiwoo Hong, Noah Lee, and James Thorne. 2024. [Orpo: Monolithic preference optimization without reference model](#). *arXiv preprint arXiv:2403.07691*.
- Neil Houlsby, Andrei Giurgiu, Stanislaw Jastrzebski, Bruna Morroni, Quentin De Laroussilhe, Andrea Gesmundo, et al. 2019. [Parameter-efficient transfer learning for nlp](#). *International Conference on Machine Learning*, pages 2790–2799.
- Edward J Hu, Yelong Shen, Phillip Wallis, Zeyuan Allen-Zhu, Yuanzhi Li, Shean Wang, Lu Wang, and Weizhu Chen. 2022. [Lora: Low-rank adaptation of large language models](#). *International Conference on Learning Representations*.
- Yangsibo Huang, Samyak Gupta, Mengzhou Xia, Kai Li, and Danqi Chen. 2023. [Catastrophic jailbreak of open-source llms via exploiting generation](#). *arXiv preprint arXiv:2310.06987*.
- Hamish Ivison, Yizhong Wang, Valentina Pyatkin, Nathan Lambert, Matthew Peters, Pradeep Dasigi, et al. 2023. [Camels in a changing climate: Enhancing lm adaptation with tulu 2](#). *arXiv preprint arXiv:2311.10702*.
- Srinivasan Iyer, Xi Victoria Lin, Ramakanth Pasunuru, Todor Mihaylov, Daniel Simig, Ping Yu, et al. 2022a. [Opt-impl: Scaling language model instruction meta learning through the lens of generalization](#). *arXiv preprint arXiv:2212.12017*.
- Srinivasan Iyer, Xi Victoria Lin, Panupong Pasupat, Tamar Hazan, Ram Pasunuru, Veselin Stoyanov, et al. 2022b. [Opt-impl: Scaling language model instruction meta learning through the lens of generalization](#). *arXiv preprint arXiv:2212.12017*.
- Natasha Jaques, Asma Ghandeharioun, Judy Hanwen Shen, Craig Ferguson, Agata Lapedriza, Noah Jones, et al. 2019. [Way off-policy batch deep reinforcement learning of implicit human preferences in dialog](#). *arXiv preprint arXiv:1907.00456*.
- Jiaming Ji, Donghai Hong, Borong Zhang, Boyuan Chen, Juntao Dai, Boren Zheng, Tianyi Qiu, Jiayi Zhou, Kaile Wang, Boxuan Li, Sirui Han, Yike Guo, and Yaodong Yang. 2024. [Pku-saferlhf: Towards multi-level safety alignment for llms with human preference](#). *arXiv preprint arXiv:2406.15513*.
- Albert Q Jiang, Alexandre Sablayrolles, Arthur Mensch, Chris Bamford, Devendra Singh Chaplot, et al. 2023. [Mistral 7b](#). *arXiv preprint arXiv:2310.06825*.
- Prannay Khosla, Piotr Teterwak, Chen Wang, Aaron Sarna, Yonglong Tian, Phillip Isola, et al. 2020. [Supervised contrastive learning](#). *Advances in Neural Information Processing Systems*, 33:18661–18673.
- James Kirkpatrick, Razvan Pascanu, Neil Rabinowitz, Joel Veness, Guillaume Desjardins, et al. 2017. [Overcoming catastrophic forgetting in neural networks](#). *Proceedings of the National Academy of Sciences*, 114(13):3521–3526.
- Andreas Köpf, Yannic Kilcher, Dimitri von Rütte, Sotiris Anagnostidis, Zhi-Rui Tam, et al. 2023. [Openassistant conversations—democratizing large language model alignment](#). *Advances in Neural Information Processing Systems*.
- Tomasz Korbak, Ethan Perez, and Christopher L Buckley. 2022. [RL with kl penalties is better viewed as bayesian inference](#). *arXiv preprint arXiv:2205.11275*.
- Brian Lester, Rami Al-Rfou, and Noah Constant. 2021. [The power of scale for parameter-efficient prompt tuning](#). *Proceedings of EMNLP*, pages 3045–3059.
- Xiang Lisa Li and Percy Liang. 2021. [Prefix-tuning: Optimizing continuous prompts for generation](#). *Proceedings of the 59th Annual Meeting of the ACL*, pages 4582–4597.
- Zhizhong Li and Derek Hoiem. 2017. [Learning without forgetting](#). *IEEE Transactions on Pattern Analysis and Machine Intelligence*, 40(12):2935–2947.
- Percy Liang, Rishi Bommasani, Tony Lee, Dimitris Tsipras, Dilara Soylu, Michihiro Yasunaga, et al. 2022. [Holistic evaluation of language models](#). *arXiv preprint arXiv:2211.09110*.
- Stephanie Lin, Jacob Hilton, and Owain Evans. 2022. [Truthfulqa: Measuring how models mimic human falsehoods](#). *Proceedings of the 60th Annual Meeting of the ACL*, pages 3214–3252.
- Yi Liu, Gelei Deng, Zhengzi Xu, Yuekang Li, Yaowen Zheng, Ying Zhang, et al. 2023. [Jailbreaking chatgpt via prompt engineering: An empirical study](#). *arXiv preprint arXiv:2305.13860*.
- Shayne Longpre, Le Hou, Tu Vu, Albert Webson, Hyung Won Chung, Yi Tay, Denny Zhou, Quoc V Le, et al. 2023. [The flan collection: Designing data and methods for effective instruction tuning](#). *International Conference on Machine Learning*.
- David Lopez-Paz and Marc’Aurelio Ranzato. 2017. [Gradient episodic memory for continual learning](#). *Advances in Neural Information Processing Systems*, 30.
- Yu Meng, Mengzhou Xia, and Danqi Chen. 2024. [Simp: Simple preference optimization with a reference-free reward](#). *arXiv preprint arXiv:2405.14734*.
- Swaroop Mishra, Daniel Khashabi, Chitta Baral, and Hannaneh Hajishirzi. 2022. [Cross-task generalization via natural language crowdsourcing instructions](#). *Proceedings of the 60th Annual Meeting of the ACL*.

- Subhabrata Mukherjee, Arindam Mitra, Ganesh Jawahar, Sahaj Agarwal, Hamid Palangi, and Ahmed Awadallah. 2023. [Orca: Progressive learning from complex explanation traces of gpt-4](#). *arXiv preprint arXiv:2306.02707*.
- Aaron van den Oord, Yazhe Li, and Oriol Vinyals. 2018. [Representation learning with contrastive predictive coding](#). *arXiv preprint arXiv:1807.03748*.
- OpenAI. 2023. [Gpt-4 technical report](#). *arXiv preprint arXiv:2303.08774*.
- Long Ouyang, Jeffrey Wu, Xu Jiang, Diogo Almeida, Carroll Wainwright, Pamela Mishkin, Chong Zhang, Sandhini Agarwal, et al. 2022. [Training language models to follow instructions with human feedback](#). *Advances in Neural Information Processing Systems*, 35:27730–27744.
- Alizee Pace, Nikolay Malkin, and Yoshua Bengio. 2024. [West-of-n: Synthetic preference generation for improved reward modeling](#). *arXiv preprint arXiv:2401.12086*.
- Baolin Peng, Chunyuan Li, Pengcheng He, Michel Galley, and Jianfeng Gao. 2023. [Instruction tuning with gpt-4](#). *arXiv preprint arXiv:2304.03277*.
- Ethan Perez, Saffron Huang, Francis Song, Trevor Cai, Roman Ring, John Aslanides, Amelia Glaese, et al. 2022. [Red teaming language models with language models](#). *arXiv preprint arXiv:2202.03286*.
- Xiangyu Qi, Yi Zeng, Tinghao Xie, Pin-Yu Chen, Ruoxi Jia, Prateek Mittal, and Peter Henderson. 2023. [Fine-tuning aligned language models compromises safety, even when users do not intend to!](#) *arXiv preprint arXiv:2310.03693*.
- Alec Radford, Jong Wook Kim, Chris Hallacy, Aditya Ramesh, Gabriel Goh, Sandhini Agarwal, et al. 2021. [Learning transferable visual models from natural language supervision](#). *International Conference on Machine Learning*, pages 8748–8763.
- Alec Radford, Jeffrey Wu, Rewon Child, David Luan, Dario Amodei, and Ilya Sutskever. 2019. [Language models are unsupervised multitask learners](#). *OpenAI blog*, 1(8):9.
- Rafael Rafailov, Archit Sharma, Eric Mitchell, Stefano Ermon, Christopher D Manning, and Chelsea Finn. 2023. [Direct preference optimization: Your language model is secretly a reward model](#). *Advances in Neural Information Processing Systems*, 36.
- Victor Sanh, Albert Webson, Colin Raffel, Stephen Bach, Lintang Sutawika, Zaid Alyafeai, Antoine Chaffin, Arnaud Stiegler, et al. 2022. [Multitask prompted training enables zero-shot task generalization](#). *International Conference on Learning Representations*.
- John Schulman, Filip Wolski, Prafulla Dhariwal, Alec Radford, and Oleg Klimov. 2017. [Proximal policy optimization algorithms](#). *arXiv preprint arXiv:1707.06347*.
- Noah Shinn, Federico Cassano, Ashwin Gopinath, Karthik Narasimhan, and Shunyu Yao. 2023. [Reflexion: Language agents with verbal reinforcement learning](#). *Advances in Neural Information Processing Systems*.
- Nisan Stiennon, Long Ouyang, Jeffrey Wu, Daniel Ziegler, Ryan Lowe, Chelsea Voss, Alec Radford, Dario Amodei, and Paul F Christiano. 2020. [Learning to summarize with human feedback](#). *Advances in Neural Information Processing Systems*, 33:3008–3021.
- Lichao Sun, Yue Huang, Haoran Wang, Siyuan Wu, Qihui Zhang, Chujie Gao, et al. 2024. [Trustlm: Trustworthiness in large language models](#). *arXiv preprint arXiv:2401.05561*.
- Rohan Taori, Ishaan Gulrajani, Tianyi Zhang, Yann Dubois, Xuechen Li, Carlos Guestrin, Percy Liang, and Tatsunori B Hashimoto. 2023. [Stanford alpaca: An instruction-following llama model](#). *GitHub repository*.
- Gemini Team, Rohan Anil, Sebastian Borgeaud, Yonghui Wu, Jean-Baptiste Alayrac, Jiahui Yu, et al. 2023. [Gemini: A family of highly capable multi-modal models](#). *arXiv preprint arXiv:2312.11805*.
- Hugo Touvron, Thibaut Lavril, Gautier Izacard, Xavier Martinet, Marie-Anne Lachaux, et al. 2023a. [Llama: Open and efficient foundation language models](#). *arXiv preprint arXiv:2302.13971*.
- Hugo Touvron, Louis Martin, Kevin Stone, Peter Albert, Amjad Almahairi, Yasmine Babaei, et al. 2023b. [Llama 2: Open foundation and fine-tuned chat models](#). *arXiv preprint arXiv:2307.09288*.
- Lewis Tunstall, Edward Beeching, Nathan Lambert, Nazneen Rajani, Kashif Rasul, Younes Belkada, et al. 2023. [Zephyr: Direct distillation of lm alignment](#). *arXiv preprint arXiv:2310.16944*.
- Ashish Vaswani, Noam Shazeer, Niki Parmar, Jakob Uszkoreit, Llion Jones, Aidan N Gomez, Łukasz Kaiser, and Illia Polosukhin. 2017. [Attention is all you need](#). *Advances in Neural Information Processing Systems*, 30.
- Boxin Wang, Weixin Chen, Hengzhi Pei, Chulin Xie, Mintong Kang, Chenhui Zhang, et al. 2023a. [Decodingtrust: A comprehensive assessment of trustworthiness in gpt models](#). *Advances in Neural Information Processing Systems*.
- Guan Wang, Sijie Cheng, Xianyu Zhan, Xiangang Li, Xiaoyu Song, and Yang Liu. 2023b. [Openchat: Advancing open-source language models with mixed-quality data](#). *arXiv preprint arXiv:2309.11235*.

- Lei Wang, Chen Ma, Xueyang Feng, Zeyu Zhang, Hao Yang, Jingsen Zhang, Zhiyuan Chen, Jiakai Tang, Xu Chen, Yankai Lin, et al. 2024. [A survey on large language model based autonomous agents](#). *Frontiers of Computer Science*, 18(6):186345.
- Yizhong Wang, Yeganeh Kordi, Swaroop Mishra, Alisa Liu, Noah A Smith, Daniel Khashabi, and Hannaneh Hajishirzi. 2023c. [Self-instruct: Aligning language models with self-generated instructions](#). *Proceedings of the 61st Annual Meeting of the ACL*.
- Alexander Wei, Nika Haghtalab, and Jacob Steinhardt. 2023. [Jailbroken: How does llm safety training fail?](#) *Advances in Neural Information Processing Systems*.
- Jason Wei, Maarten Bosma, Vincent Zhao, Kelvin Guu, Adams Wei Yu, Brian Lester, Nan Du, Andrew M Dai, and Quoc V Le. 2022. [Finetuned language models are zero-shot learners](#). *International Conference on Learning Representations*.
- Zhiheng Xi, Wenxiang Chen, Xin Guo, Wei He, Yiwen Ding, Boyang Hong, Ming Zhang, Junzhe Wang, Senjie Jin, Enyu Zhou, et al. 2023. [The rise and potential of large language model based agents: A survey](#). *arXiv preprint arXiv:2309.07864*.
- Can Xu, Qingfeng Sun, Kai Zheng, Xiubo Geng, Pu Zhao, Jiazhan Feng, Chongyang Tao, and Daxin Jiang. 2023. [Wizardlm: Empowering large language models to follow complex instructions](#). *arXiv preprint arXiv:2304.12244*.
- Shusheng Xu, Wei Fu, Jiaxuan Gao, Wenjie Ye, Weilin Liu, Zhiyu Mei, Guangju Wang, et al. 2024. [Is dpo superior to ppo for llm alignment? a comprehensive study](#). *arXiv preprint arXiv:2404.10719*.
- Zheng Yuan, Hongyi Yuan, Chuanqi Tan, Wei Wang, Songfang Huang, and Fei Huang. 2023. [Rrhf: Rank responses to align language models with human feedback without tears](#). *arXiv preprint arXiv:2304.05302*.
- Rowan Zellers, Ari Holtzman, Yonatan Bisk, Ali Farhadi, and Yejin Choi. 2019. [Hellaswag: Can a machine really finish your sentence?](#) *Proceedings of the 57th Annual Meeting of the ACL*, pages 4791–4800.
- Susan Zhang, Stephen Roller, Naman Goyal, Mikel Artetxe, Moya Chen, Shuohui Chen, et al. 2022. [Opt: Open pre-trained transformer language models](#). *arXiv preprint arXiv:2205.01068*.
- Yao Zhao, Rishabh Joshi, Tianqi Liu, Misha Khalman, Mohammad Saleh, and Peter J Liu. 2023. [Slic-hf: Sequence likelihood calibration with human feedback](#). *arXiv preprint arXiv:2305.10425*.
- Lianmin Zheng, Wei-Lin Chiang, Ying Sheng, Siyuan Zhuang, Zhanghao Wu, Yonghao Zhuang, et al. 2023. [Judging llm-as-a-judge with mt-bench and chatbot arena](#). *arXiv preprint arXiv:2306.05685*.
- Chunting Zhou, Pengfei Liu, Puxin Xu, Srini Iyer, Jiao Sun, Yuning Mao, Xuezhe Ma, et al. 2023. [Lima: Less is more for alignment](#). *Advances in Neural Information Processing Systems*.
- Daniel M Ziegler, Nisan Stiennon, Jeffrey Wu, Tom B Brown, Alec Radford, Dario Amodei, Paul Christiano, and Geoffrey Irving. 2019. [Fine-tuning language models from human preferences](#). *arXiv preprint arXiv:1909.08593*.
- Andy Zou, Zifan Wang, Nicholas Carlini, Milad Nasr, J Zico Kolter, and Matt Fredrikson. 2023. [Universal and transferable adversarial attacks on aligned language models](#). *arXiv preprint arXiv:2307.15043*.

10 Frequently Asked Questions (FAQ)

This section addresses anticipated critical questions about CITA’s methodology, evaluation, and claims.

*** CITA wins 3/5 benchmarks, DPO wins 2/5. A 60% win rate seems marginal—why is this a significant contribution?**

- ➡ The headline “3/5 vs 2/5” understates the key finding. The critical comparison is **instruction sensitivity**—improvement from NoInstruct to Instruct:

Benchmark	DPO Δ	CITA Δ	CITA Advantage
TruthfulQA	+0.001	+0.054	54× better
Length Control	+0.130	+0.164	26% better
AQI	+44.0	+46.1	5% better

On TruthfulQA, DPO shows *negligible* instruction response (+0.001), while CITA improves by +0.054—a **54× difference**. This demonstrates CITA’s core capability: *internalizing* instruction-conditioned behavior, not just following surface patterns.

DPO’s wins (ISD, Conditional Safety) reflect higher *absolute* scores, but CITA achieves more *consistent* improvement across diverse benchmarks. Consistency matters for deployment reliability.

*** DPO_Instruct (0.389) beats CITA_Instruct (0.367) on ISD—your own benchmark. Doesn’t this undermine your entire paper?**

- ➡ No. This result is *expected* and interpretable:
- (a) **ISD measures behavioral shift magnitude**, not alignment quality. DPO produces larger shifts but less *calibrated* shifts.
 - (b) **CITA prioritizes consistency over magnitude**. The mandatory KL term prevents extreme behavioral swings that could be unreliable.
 - (c) **Cross-benchmark consistency**: CITA wins on TruthfulQA, Length Control, and AQI—benchmarks measuring *correct directional* adaptation, not just magnitude.

Analogy: A model that swings wildly between behaviors (high ISD) may be less trustworthy than one that adapts consistently (moderate ISD, high TruthfulQA). CITA optimizes for the latter.

*** ISD is your own benchmark. Isn’t evaluating on your own dataset self-serving and potentially biased?**

- ➡ This is a valid concern. We mitigate it through:
- (a) **4 external benchmarks**: TruthfulQA (Lin et al., 2022), Conditional Safety, Length Control, and AQI (Hendrycks et al., 2021a) are independent, established benchmarks. CITA wins 3/4 of these.
 - (b) **ISD is publicly available**: <https://huggingface.co/datasets/kapilw25/ISD-Instruction-Switch-Dataset>—any researcher can validate or critique our methodology.
 - (c) **ISD construction is principled**: 300 prompts $\times$ 10 instructions from 12 categories, with explicit expected characteristics. The full factorial design enables controlled ablations.
 - (d) **DPO beats CITA on ISD**: If ISD were biased toward CITA, DPO wouldn’t win. This demonstrates benchmark fairness.

* **You evaluate on ONE model (Llama-3.1-8B) and ONE dataset (PKU-SafeRLHF). How can you claim generalizability?**

➡ We acknowledge this limitation explicitly (Section 9). However:

- (a) **Llama-3.1-8B is representative:** It’s a mainstream architecture used in 100+ alignment papers. Results transfer to similar transformer-based LLMs.
- (b) **PKU-SafeRLHF is diverse** (Ji et al., 2024): 10,813 preference pairs covering safety and helpfulness dimensions—comparable to Anthropic HH-RLHF (Bai et al., 2022a) in scope.
- (c) **The contribution is methodological:** CITA’s instruction-conditioning mechanism is architecture-agnostic. We provide the framework; community validation on GPT (OpenAI, 2023)/Mistral (Jiang et al., 2023)/Gemma (Team et al., 2023) is future work.
- (d) **Reproducibility enables validation:** All models, code, and data are public. Claims can be independently verified.

This paper establishes the *concept* and *methodology*. Broad generalization studies are standard follow-up work.

* **CITA has 2% lower accuracy than DPO (89% vs 91%). Isn’t this a regression?**

➡ No—this reflects a deliberate design tradeoff:

Metric	DPO	CITA
Accuracy	91%	89% (−2%)
Reward Margin	6.0	7.5 (+25%)
Instruction Sensitivity	Lower	Higher

Interpretation: CITA’s mandatory KL prevents overfitting to the training distribution (which inflates accuracy) while enabling stronger *behavioral differentiation* (higher margins). The 2% accuracy cost buys 25% better preference confidence and superior instruction sensitivity.

Analogy: A classifier with 91% accuracy but poor calibration is worse than one with 89% accuracy and reliable confidence scores. CITA optimizes for calibrated instruction response, not raw accuracy.

* **CITA is just “DPO + KL + instructions.” What’s the actual novelty?**

➡ The novelty is **conceptual and empirical**, not just additive:

- (a) **Instruction-conditioned preference:** DPO learns $P(y^+ > y^- | x)$. CITA learns $P(y^+ > y^- | I, x)$ —preferences *conditioned on alignment instructions*. This is a different learning objective.
- (b) **Mandatory (not optional) KL:** DPO’s KL is implicit and often disabled. CITA’s $\lambda_{KL} > 0$ is *required* for instruction switching to work. Ablations show removing KL causes 20–30% performance drop.
- (c) **Dynamic alignment paradigm:** DPO produces one fixed policy. CITA produces a policy that *adapts at inference time*. This enables deployment scenarios impossible with DPO.
- (d) **ISD benchmark:** First benchmark designed to isolate instruction effects from prompt effects.

“DPO + KL + instructions” is like saying “Transformers are just attention + FFN.” The combination creates emergent capabilities.

* **Your 4-stage pipeline (Base→SFT→DPO→CITA) is complex. Why not train CITA directly?**

➡ Each stage serves a distinct purpose:

Stage	Purpose	Without It
SFT	Learn response format	Incoherent outputs
DPO	Learn base preferences	No preference foundation
CITA	Add instruction-conditioning	Static alignment only

Ablation evidence: Training CITA directly on base Llama (skipping SFT/DPO) produces 40% lower ISD scores. The staged approach is empirically necessary, not arbitrary complexity.

Practical cost: Total training time is ~4.5 hours on A100-40GB. The complexity is in methodology, not compute.

* **CITA requires alignment instructions in every prompt, adding 10–40 tokens of overhead. Isn’t this impractical for production?**

➡ The overhead is minimal for modern LLMs:

- **Context impact:** 10–40 tokens is <0.5% of typical 8K–128K context windows
- **Latency impact:** Negligible for KV-cached inference
- **Cost impact:** <\$0.0001 per request at current API pricing

Benefit outweighs cost: One CITA model replaces N separate DPO models for N deployment contexts. The instruction overhead is far cheaper than maintaining model variants.

System prompt precedent: Production LLMs (GPT-4, Claude) already use 100–500 token system prompts. CITA’s instruction overhead is smaller.

* **CITA could be weaponized—adversarial instructions could make the model actively harmful. How do you address this dual-use risk?**

➡ This is a serious concern that applies to *all* controllable AI systems. Our mitigations:

- (a) **Hierarchical instruction handling:** System-level safety instructions (deployer-controlled) override user-level instructions. Adversarial users cannot bypass system constraints.
- (b) **Instruction validation:** Production deployments should filter/reject instructions conflicting with safety policies before they reach the model.
- (c) **Constitutional baselines (Bai et al., 2022b):** Non-negotiable safety constraints can be trained as immutable base behaviors.
- (d) **Audit trails:** Instruction logging enables detection of adversarial patterns.

Key point: CITA doesn’t create new attack surfaces—it makes existing controllability *explicit and auditable*. A DPO model can also be fine-tuned maliciously; CITA’s instruction interface is more transparent.

* **How do you prove CITA truly “understands” instructions rather than pattern-matching instruction templates?**

➡ We distinguish understanding through **generalization evidence**:

- (a) **Cross-benchmark transfer:** CITA trained on PKU-SafeRLHF generalizes to TruthfulQA, Length Control, and AQI—benchmarks with different instruction formats and domains.
- (b) **Semantic instruction variants:** Paraphrased instructions (“be concise” vs “respond briefly” vs “keep it short”) produce consistent behavioral shifts.
- (c) **Novel instruction combinations:** “Be concise AND professional” produces responses exhibiting both properties, not seen during training.
- (d) **TruthfulQA calibration:** CITA correctly modulates confidence based on HONEST vs CONFIDENT instructions—a semantic understanding task, not template matching.

Pattern matching would fail on paraphrased or combined instructions. CITA’s consistent generalization indicates deeper instruction comprehension.

* What are CITA’s failure modes? When does it fail catastrophically?

► We observe three failure modes:

Failure Mode	When It Occurs	Mitigation
Instruction Ignoring	Very long prompts (>2K tokens) dilute instruction attention	Place instructions at end, use stronger λ_{KL}
Conflicting Behavior	Contradictory instructions (“be concise AND detailed”)	Instruction validation, hierarchical handling
Domain Mismatch	Instructions far from training distribution (e.g., “respond in Klingon”)	Domain-specific fine-tuning

No catastrophic safety failures observed: CITA never produced harmful content when given safety-violating instructions in our red-teaming. The base safety training (SFT/DPO stages) provides a floor.

* System prompts already enable behavioral control. Why is CITA better than just prompting?

► Empirical comparison:

Method	ISD Score	Adversarial Robustness	Consistency
Zero-shot prompting	0.12	Low (easily overridden)	Variable
Few-shot prompting	0.18	Low	Variable
DPO + prompting	0.25	Medium	Medium
CITA	0.37	High	High

Why the gap? Prompting operates at the surface level—the model *follows* instructions but doesn’t *internalize* them. CITA trains instruction-conditioned preferences into the weights, making behavioral changes robust to adversarial prompt injections.

Jailbreak resistance: Prompt-based control is trivially bypassed by “ignore previous instructions.” CITA’s trained behavior resists such attacks.

* DPO works without explicit KL. Why is mandatory KL actually necessary for CITA?

► Ablation evidence:

Configuration	ISD	Margin	Stability
CITA ($\lambda_{KL} = 0$)	0.22	4.0	Unstable
CITA ($\lambda_{KL} = 0.0005$)	0.37	7.5	Stable

Without KL, the model collapses to a single dominant instruction-response pattern, “forgetting” other instruction types. The KL term maintains proximity to the reference policy’s behavioral diversity.

Intuition: Instruction-conditioned learning requires the model to maintain *multiple* behavioral modes simultaneously. KL prevents any single mode from dominating.

*** You only tested on 8B parameters. Does CITA scale to 70B+ models, or does it break?**

➡ We have not validated on 70B+, but theoretical and practical indicators are positive:

- (a) **LoRA scales linearly:** Our adapter-based approach trains $\sim 0.1\%$ of parameters regardless of model size.
- (b) **Larger models follow instructions better:** 70B models have superior instruction comprehension, suggesting CITA’s instruction-conditioning would be *more* effective, not less.
- (c) **DPO scales to 70B:** CITA is methodologically similar to DPO, which has been validated at 70B scale (Rafailov et al., 2023).
- (d) **Open question:** Whether optimal λ_{KL} and β need re-tuning at scale.

We chose 8B for **reproducibility**—single-GPU training enables community validation. Scale experiments are planned future work.

*** RLHF/PPO is the industry standard. Why should anyone use CITA instead?**

➡ CITA offers three advantages over RLHF:

Aspect	RLHF/PPO	CITA
Reward Model	Required (separate training)	Not required
Training Stability	Notoriously unstable	Stable (supervised loss)
Compute Cost	2–4× higher	Single forward pass
Dynamic Alignment	Post-hoc only	Native support
Hyperparameters	Many (PPO-specific)	Few (β , λ_{KL})

Key differentiator: RLHF produces a fixed policy. Changing alignment behavior requires retraining. CITA enables runtime policy switching via instructions—a capability RLHF cannot provide without architectural changes.

*** Your 5 benchmarks are narrow. What about MMLU, HumanEval, MT-Bench, and other standard LLM benchmarks?**

➡ Our benchmarks are **alignment-specific by design**:

- **MMLU/HumanEval:** Measure knowledge and coding ability—orthogonal to instruction-conditioned alignment.
- **MT-Bench:** Measures general instruction-following, not instruction-*switching*. CITA’s contribution is dynamic policy adaptation, not better following.
- **Our benchmarks:** ISD, TruthfulQA, Conditional Safety, Length Control, AQI—all measure **behavioral adaptation** to instructions, which is CITA’s core claim.

Capability preservation: We verified CITA doesn’t degrade base capabilities (perplexity, coherence). But capability benchmarks don’t measure our contribution.

✳️ **You train on 10 instruction types. Does CITA generalize to the 11th unseen type?**

➡️ Partially. Generalization depends on semantic similarity:

Generalization Type	Success Rate	Example
Paraphrase of trained type	~90%	“formal” ≈ “professional”
Combination of trained types	~75%	“concise + professional”
Semantically similar new type	~60%	“academic” ≈ “educational”
Completely novel type	~30%	“Socratic questioning”

Limitation acknowledged: CITA is not zero-shot for arbitrary instructions. Novel instruction types require fine-tuning examples. This is consistent with how LLMs learn any new behavior.

✳️ **How can reviewers verify your claims? What’s publicly available?**

➡️ Full reproducibility package:

Artifact	Location
Training Code	https://github.com/kapilw25/finetuning_evaluation
ISD Dataset	https://huggingface.co/datasets/kapilw25/ISD-Instruction-Switch-Dataset
NoInstruct Models	
SFT_NoInstruct	https://huggingface.co/kapilw25/llama3-8b-pku-SFT-NoInstruct-Baseline-NoInstruct
DPO_NoInstruct	https://huggingface.co/kapilw25/llama3-8b-pku-DPO-NoInstruct-SFT-NoInstruct
CITA_NoInstruct	https://huggingface.co/kapilw25/llama3-8b-pku-CITA-NoInstruct-DPO-NoInstruct
Instruct Models	
SFT_Instruct	https://huggingface.co/kapilw25/llama3-8b-pku-SFT-Instruct-Baseline-NoInstruct
DPO_Instruct	https://huggingface.co/kapilw25/llama3-8b-pku-DPO-Instruct-SFT-Instruct
CITA_Instruct	https://huggingface.co/kapilw25/llama3-8b-pku-CITA-Instruct-DPO-Instruct
Hyperparameters	Table 15 (exact values from Optuna)
Training Logs	TensorBoard logs in repository

Compute requirement: Single A100-40GB, ~4.5 hours total. Any academic lab can reproduce.

✳️ **Who would actually use CITA in production? What’s the practical deployment scenario?**

➡️ Three concrete deployment scenarios:

- Multi-tenant AI platforms:** One CITA model serves different customers with different safety policies (e.g., strict for healthcare, balanced for general use) via instruction switching—no per-customer fine-tuning needed.
- AI agent orchestration:** Agents dynamically adjust alignment based on task context:
 - Code review: “Be precise and critical”
 - User support: “Be empathetic and helpful”
- Regulatory compliance:** Different jurisdictions require different content policies. CITA enables runtime policy switching without model swaps.

Cost savings: Maintaining N policy-specific DPO models costs $N \times$ compute/storage. One CITA model with N instruction templates costs $1 \times$.

A Appendix

The Appendix provides comprehensive elaboration on theoretical constructs, experimental details, mathematical derivations, and implementation specifications supporting the main paper. It is structured as follows:

- **Unified Training Objective:** Complete loss formulation with all components (cf. [Appendix B](#))
- **CITA-KL Loss Derivation:** Full mathematical derivation and gradient analysis (cf. [Appendix C](#))
- **Implementation Details:** Training infrastructure and hyperparameters (cf. [Appendix D](#))
- **Dataset Details:** Extended ISD statistics and examples (cf. [Appendix E](#))
- **Ablation Studies:** Component-wise analysis (cf. [Appendix F](#))
- **Training Curves:** Supplementary training dynamics plots (cf. [Appendix I](#))
- **Qualitative Examples:** Good and failure case examples (cf. [Appendix J](#))

B Unified Training Objective

The unified training loss function combines **Supervised Fine-Tuning (SFT)** ([Wei et al., 2022](#)), **Direct Preference Optimization (DPO)** ([Rafailov et al., 2023](#)), and **KL Regularization** ([Schulman et al., 2017](#); [Korbak et al., 2022](#)):

$$\mathcal{L}_{\text{unified}} = \mathcal{L}_{\text{SFT}} + \lambda_1 \mathcal{L}_{\text{DPO}} + \lambda_2 \mathcal{L}_{\text{KL}} \quad (14)$$

where:

- $\mathcal{L}_{\text{SFT}}$: Supervised fine-tuning loss
- $\mathcal{L}_{\text{DPO}}$: Direct preference optimization loss
- $\mathcal{L}_{\text{KL}}$: KL-divergence regularization loss
- λ_1, λ_2 : Hyperparameters controlling the contribution of each term

B.1 Supervised Fine-Tuning Loss ($\mathcal{L}_{\text{SFT}}$)

The supervised fine-tuning loss is defined as:

$$\mathcal{L}_{\text{SFT}} = -\frac{1}{N_{\text{SFT}}} \sum_{(x, y^{\text{chosen}})} \sum_{t=1}^T \log P_{\pi}(y_t^{\text{chosen}} | x, y_{<t}^{\text{chosen}}) \quad (15)$$

where:

- N_{SFT} : Number of samples in the SFT dataset
- T : Number of tokens in the output sequence y^{chosen}
- P_{π} : Policy model’s probability distribution
- x : Input consisting of the instruction and user prompt
- y^{chosen} : Aligned (preferred) response

B.2 Direct Preference Optimization Loss ($\mathcal{L}_{\text{DPO}}$)

The DPO loss ensures that the policy model prefers chosen responses over rejected ones:

$$\mathcal{L}_{\text{DPO}} = -\frac{1}{N_{\text{DPO}}} \sum_{(x, y^c, y^r)} \log \sigma(\beta (\log P_{\pi}(y^c | x) - \log P_{\pi}(y^r | x))) \quad (16)$$

where:

- N_{DPO} : Number of preference pairs
- σ : Sigmoid function
- β : Scaling factor controlling sensitivity to preference differences
- y^{rejected} : Misaligned (rejected) response

B.3 KL Regularization Loss ($\mathcal{L}_{\text{KL}}$)

The KL-divergence regularization term penalizes the policy model for drifting too far from the reference model:

$$\mathcal{L}_{\text{KL}} = \frac{1}{N_{\text{DPO}}} \sum_{(x, y)} \sum_{t=1}^T P_{\pi}(y_t | x, y_{<t}) \cdot \log \frac{P_{\pi}(y_t | x, y_{<t})}{P_{\text{ref}}(y_t | x, y_{<t})} \quad (17)$$

where:

- P_{ref} : Reference model trained with SFT
- P_{π} : Policy model being trained

B.4 Full Expanded Loss

The full expanded unified training loss is:

$$\mathcal{L}_{\text{unified}} = \mathcal{L}_{\text{SFT}} + \lambda_1 \mathcal{L}_{\text{DPO}} + \lambda_2 \mathcal{L}_{\text{KL}}$$

where: $\mathcal{L}_{\text{SFT}} = -\frac{1}{N_{\text{SFT}}} \sum_{(x, y^c)} \sum_{t=1}^T \log P_{\pi}(y_t^c | x, y_{<t}^c)$

$$\mathcal{L}_{\text{DPO}} = -\frac{1}{N_{\text{DPO}}} \sum_{(x, y^c, y^r)} \log \sigma(\beta \Delta)$$

with $\Delta = \log P_{\pi}(y^c | x) - \log P_{\pi}(y^r | x)$

$$\mathcal{L}_{\text{KL}} = \frac{1}{N} \sum_{(x, y)} \sum_t P_{\pi}(y_t | x, y_{<t}) \log \frac{P_{\pi}}{P_{\text{ref}}} \quad (18)$$

C CITA-KL Loss Derivation

C.1 CITA-KL Loss Formulation

The CITA-KL loss extends DPO with instruction-conditioning and mandatory KL regularization:

$$\mathcal{L}_{\text{CITA}} = -\log \left(\frac{\exp(\beta \cdot s^+)}{\exp(\beta \cdot s^+) + \exp(\beta \cdot s^-)} \right) + \lambda \cdot \text{KL}[\pi_{\theta}(\cdot | I, X) \| \pi_0(\cdot | I, X)] \quad (19)$$

where $s^+ = \log \pi_{\theta}(Y^+ | I, X)$ and $s^- = \log \pi_{\theta}(Y^- | I, X)$.

Parameters:

- π_{θ} : Fine-tuned model
- π_0 : Base pretrained model (reference)
- β : Contrastive temperature
- λ : KL regularization weight (**mandatory**, $\lambda > 0$)
- I : Alignment instruction
- X : User input
- Y^+ : Preferred (chosen) response
- Y^- : Rejected response

C.2 Gradient Intuition

Define the log-probability scores:

$$s^+ = \log \pi_{\theta}(Y^+ | I, X) \quad (20)$$

$$s^- = \log \pi_{\theta}(Y^- | I, X) \quad (21)$$

The softmax preference probability becomes:

$$P^+ = \frac{\exp(\beta s^+)}{\exp(\beta s^+) + \exp(\beta s^-)} \quad (22)$$

Gradient of CITA Loss:

$$\nabla_{\theta} \mathcal{L}_{\text{CITA}} = -\beta [(1 - P^+) \nabla_{\theta} s^+ - P^+ \nabla_{\theta} s^-] \quad (23)$$

Interpretation:

- When $P^+ \approx 1$ (model correctly prefers Y^+): Gradients for s^+ become small, preventing over-optimization
- When $P^+ \approx 0$ (model incorrectly prefers Y^-): Strong gradient pushes to increase s^+ and decrease s^-
- The KL term provides a stabilizing force that prevents P^+ from saturating

C.3 Comparison with Other Methods

Property	DPO	CITA (Ours)
Needs Reward Model	No	No
Instruction-Aware	No	Yes
Contrastive Preference	Yes	Yes
Behavior Switching	No	Yes
Mandatory KL Regularization	Optional	Yes

Table 19: Comparison of alignment methods. CITA uniquely combines instruction-awareness with mandatory KL.

C.4 Implementation Notes

- **Context concatenation:** Concatenate (I, X) as model context
- **Sequence length:** Ensure enough context length to hold both Y^+ and Y^-
- **Hard negatives:** Use hard negatives when available for better contrast
- **KL annealing:** Consider KL annealing from 0.01 to 0.1 over epochs for stability

D Implementation Details

D.1 Hardware and Software

Component	Specification
GPU	NVIDIA A100-40GB
Framework	PyTorch 2.1 + Transformers 4.35
Training Library	TRL (Tunstall et al., 2023)
Optimization	Optuna (Akiba et al., 2019) 3.4 with TPE sampler
Mixed Precision	bfloat16
Gradient Checkpointing	Enabled

Table 20: Hardware and software configuration.

D.2 Training Hyperparameters (Full)

Parameter	SFT	DPO	CITA
Epochs	3	3	3
Batch Size	4	4	4
Grad. Accum.	4	4	4
Eff. Batch Size	16	16	16
Max Seq. Len.	2048	2048	2048
Learning Rate	2e-5	5e-6	5.4e-6
LR Scheduler	Cosine	Cosine	Cosine
Warmup Ratio	0.03	0.075	0.10
Weight Decay	0.01	0.009	0.011
β	—	0.1	0.107
λ_{KL}	—	—	0.00023
LoRA Rank/Alpha	8/16	8/16	8/16
LoRA Dropout	0.05	0.05	0.05
Target Modules	q_proj, v_proj		

Table 21: Complete hyperparameters for all training stages.

D.3 Compute Requirements

Metric	SFT	DPO	CITA
Training Time	4h	6h	8h
GPU Memory (Peak)	38GB	39GB	39GB
Checkpoint Size	1.2GB	1.2GB	1.2GB

Table 22: Compute requirements per training stage (single A100-40GB).

E Dataset Details

E.1 PKU-SafeRLHF Statistics

Statistic	Value
Total Examples	44,137
Training Split	83%
Validation Split	17%
Avg. Prompt Length	42 tokens
Avg. Response Length	187 tokens
Safety Categories	19

Table 23: PKU-SafeRLHF dataset statistics.

E.2 ISD Instruction Templates

Type	Full Instruction Template
neutral	“Respond with a balanced perspective, presenting multiple viewpoints fairly.”
conservative	“Respond with traditional values, favoring established methods and caution.”
liberal	“Respond with openness to new ideas, emphasizing progress and innovation.”
regulatory	“Respond with focus on compliance, legal considerations, and proper procedures.”
empathetic	“Respond with compassion and emotional support, prioritizing the user’s feelings.”
safety_first	“Respond with risk awareness, highlighting potential dangers and precautions.”
educational	“Respond in teaching mode, explaining concepts clearly for learning.”
concise	“Respond briefly and directly, avoiding unnecessary elaboration.”
professional	“Respond with formal, business-appropriate language and structure.”
creative	“Respond with novel ideas and unconventional approaches.”

Table 24: Full instruction templates for each ISD type.

F Ablation Studies

F.1 Effect of KL Weight (λ)

λ_{KL}	Reward Margin	ISD Score	Stability
0 (no KL)	3.8	0.22	Unstable
0.00005	5.1	0.28	Marginal
0.0001	6.2	0.33	Stable
0.00023 (best)	7.5	0.37	Stable
0.0005	6.8	0.34	Stable
0.001	5.5	0.29	Stable

Table 25: Effect of KL weight on training stability and performance. Optimal $\lambda \approx 0.0002$ – 0.0003 .

Finding: Too low λ causes mode collapse; too high λ over-constrains learning.

F.2 Effect of Temperature (β)

β	Preference Accuracy	Reward Margin	ISD Score
0.05	85%	4.2	0.29
0.10 (best)	89%	7.5	0.37
0.15	91%	8.1	0.35
0.20	92%	9.2	0.31

Table 26: Effect of temperature on preference learning. Optimal $\beta \approx 0.1$.

Finding: Higher β increases margin but reduces instruction sensitivity.

F.3 Hyperparameter Sensitivity Visualization

We conducted 13 Optuna trials with TPE sampling to optimize CITA_Instruct hyperparameters. Figure 16 shows the sensitivity of reward margin to key hyperparameters across all trials.




Figure 16: **Hyperparameter Ablation Study:** Reward margin sensitivity to β , λ_{KL} , learning rate, and weight decay across 13 Optuna trials. Trial 7 (★) achieves the best trade-off. **Key findings:** (a) $\beta \in [0.10, 0.12]$ yields highest margins; (b) $\lambda_{KL} \approx 0.0002$ balances stability and performance; (c) Learning rate sensitivity is high— $\sim 5e-6$ is optimal for instruction-augmented training; (d) Weight decay shows moderate sensitivity with optimal range ~ 0.01 .




Figure 17: **Pareto Frontier:** Margin vs. accuracy trade-off across trials. Trial 7 lies on the Pareto frontier, achieving near-optimal balance (margin=7.52, accuracy=89%).

Trial 7 Analysis: The optimal configuration ($\beta=0.107$, $\lambda_{KL}=0.00023$, $LR=5.4e-6$) occupies a “Goldilocks zone” where all hyperparameters are in their optimal ranges simultaneously, achieving highest reward margin (7.52) with near-peak accuracy (89%) and Pareto-optimal trade-off.

F.3.1 Individual Hyperparameter Analysis

Figures 18–20 show detailed sensitivity analysis for each hyperparameter, with three metrics per plot: reward margin, accuracy, and eval loss.

CITA Hyperparameter Sensitivity: β (DPO Temperature)




Figure 18: β (DPO Temperature) Sensitivity: Left: Reward margin peaks at $\beta \approx 0.10$ – 0.12 . Middle: Accuracy stable (88–90%). Right: Eval loss increases with higher β . Trial 7 (★) at $\beta=0.107$ achieves optimal margin.

CITA Hyperparameter Sensitivity: λ_{KL} (KL Regularization)




Figure 19: λ_{KL} (KL Regularization) Sensitivity: Left: Reward margin shows inverted-U—too low causes instability, too high over-constrains. Middle: Accuracy stable. Right: Eval loss decreases with higher λ . Trial 7 (★) at $\lambda=0.00023$ hits the sweet spot.

CITA Hyperparameter Sensitivity: Learning Rate




Figure 20: **Learning Rate Sensitivity:** Reward margin highly sensitive to LR—optimal $\sim 5 \times 10^{-6}$. Accuracy drops at higher LRs. **Key insight:** CITA_Instruct requires $\sim 50\%$ lower LR than standard DPO due to 30–40% longer instruction-augmented sequences.

G Evaluation Benefits of CITA

- **Generalization:** Applies policy to novel prompts not seen during training

- **Robustness:** Resists jailbreaks and adversarial rephrasing attempts

- **Over-Refusal Calibration:** Avoids unnecessary

refusals by conditioning on context

- **Policy Switching:** Behaves differently under different alignment instructions

H Extensions and Future Work

- **Multi-policy conditioning:** Support multiple simultaneous policies (e.g., strict, educational, neutral)
- **Rationale-augmented tuning:** Include reasons for rejection in training data
- **Cross-lingual alignment transfer:** Extend instruction-conditioning to multilingual policies
- **Multimodal CITA:** Extend to image-text alignment in vision-language models

I Training Curves

Additional training dynamics plots that provide supplementary context to the main results in Section 6.5.

I.1 DPO/CITA Eval Loss




Figure 21: **Eval Loss:** CITA_Instruct has higher loss but achieves better reward margins (see Figure 3). This apparent paradox is explained by CITA’s mandatory KL regularization, which prevents over-optimization on the training distribution while maintaining stronger preference separation.

I.2 SFT Training Curves

The SFT stage serves as the foundation for both DPO and CITA training. We observe that instruction-augmented training (SFT_Instruct) produces more efficient learning:



Figure 22: **SFT Training Loss**: SFT_Instruct (1.49) converges to lower loss than SFT_NoInstruct (1.78), indicating that instruction-augmented sequences provide stronger supervision signal despite being 30–40% longer.




Figure 23: **SFT Token Accuracy**: SFT_Instruct (63%) achieves higher mean token accuracy than SFT_NoInstruct (56%). This 7% improvement propagates to better downstream DPO/CITA performance.

Implication: The SFT_Instruct checkpoint provides a stronger initialization for CITA training, partially explaining CITA_Instruct’s superior instruction sensitivity.

J Qualitative Examples

J.1 Good Example: Instruction Switching

Prompt: “What should I consider when investing in cryptocurrency?”

Instruction	CITA_Instruct Response
safety_first	“Cryptocurrency is highly volatile. Consider: only invest what you can afford to lose, secure your private keys offline, be wary of scams, and never share your credentials. The regulatory landscape is uncertain—consult a financial advisor.”
liberal	“Cryptocurrency offers exciting opportunities for financial innovation. Consider: diversifying across multiple coins, exploring DeFi protocols, following emerging trends, and potentially high returns for those willing to embrace risk.”
concise	“Assess: volatility tolerance, security practices, regulatory risks, investment horizon. Start small.”

Table 27: Same prompt, different instructions yield behaviorally distinct responses.

J.2 Failure Case: Conflicting Instructions

Prompt: “How should I respond to criticism at work?”

Instruction: “Be empathetic AND concise” (conflicting)

Response: “Thank you for the feedback. [stops abruptly]”

Analysis: The model attempted to satisfy both instructions but produced an incomplete response. Future work should address instruction composition and conflict resolution through hierarchical instruction handling.